

Logistic_regression_train_val_test_mnist_keras

March 7, 2018

1 Logistic Regression on MNIST

```
In [1]: import six.moves.cPickle as pickle
import numpy as np
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.datasets import mnist
from keras.utils import np_utils
from keras import optimizers
```

Using TensorFlow backend.

```
In [2]: def build_logistic_model(input_dim, output_dim):
    model = Sequential()
    model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
    return model
```

```
In [3]: batch_size = 128
nb_classes = 10
nb_epoch = 15
input_dim = 784
```

```
In [4]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
print X_train.shape
print y_train.shape
```

```
(60000, 28, 28)
(60000,)
```

```
In [5]: #converting each image into a single row vector
X_train = X_train.reshape(60000, input_dim)
X_test = X_test.reshape(10000, input_dim)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255
X_test /= 255
```

```

In [6]: # convert class vectors to binary class matrices
        Y_train = np_utils.to_categorical(y_train, nb_classes)

        #we stack the the images and there corresponding labels together.

        X_total_train=np.hstack((X_train,Y_train))
        print X_total_train.shape

        # We make sure that the data is shuffled properly so that we have a uniformly distributed
        np.random.shuffle(X_total_train)

        #Again split the data into image vectors and their corresponding labels
        X_train = X_total_train [:60000,:784]
        Y_train = X_total_train [:60000,784:794]
        print X_train.shape
        print Y_train.shape

(60000, 794)
(60000, 784)
(60000, 10)

In [7]: #Dividing the training set (X_train) of 60,000 images into a training set(X_train_t) of
        #and a validation set(X_train_v) of 10,000
        X_train_t = X_train[: 50000,:]
        X_train_v = X_train[ 50000:60000,:]
        print X_train_t.shape
        print X_train_v.shape
        print(X_train_t.shape[0], 'train samples')
        print(X_train_v.shape[0], 'validation samples')

(50000, 784)
(10000, 784)
(50000, 'train samples')
(10000, 'validation samples')

In [8]: # The corresponding labels are also diivded into training and validation set
        Y_train_t = Y_train[: 50000,:]
        Y_train_v = Y_train[50000:60000,:]
        print Y_train_t.shape
        print Y_train_v.shape

        # convert class vectors to binary class matrices for test labels
        Y_test = np_utils.to_categorical(y_test, nb_classes)
        print(Y_test.shape[0], 'Test samples')

(50000, 10)
(10000, 10)

```

```
(10000, 'Test samples')
```

```
In [9]: model = build_logistic_model(input_dim, nb_classes)
```

```
model.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
=====
dense_1 (Dense)              (None, 10)            7850
=====
Total params: 7,850
Trainable params: 7,850
Non-trainable params: 0
-----
```

```
In [10]: # compile the model
```

```
         #setting the learning rate to be 0.01
```

```
sgd = optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)
```

```
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
         #using training and validation data for 15 epochs
```

```
history = model.fit(X_train_t, Y_train_t,
                    batch_size=batch_size, nb_epoch=nb_epoch, shuffle=True,
                    verbose=1, validation_data=(X_train_v, Y_train_v))
```

```
         #evaluating the model on test set
```

```
score = model.evaluate(X_test, Y_test, verbose=0)
```

```
/usr/local/lib/python2.7/dist-packages/keras/models.py:942: UserWarning: The `nb_epoch` argument
  warnings.warn('The `nb_epoch` argument in `fit` '
```

```
Train on 50000 samples, validate on 10000 samples
```

```
Epoch 1/15
```

```
50000/50000 [=====] - 1s 21us/step - loss: 1.3573 - acc: 0.6646 - val_loss: 1.3573
```

```
Epoch 2/15
```

```
50000/50000 [=====] - 1s 16us/step - loss: 0.7720 - acc: 0.8307 - val_loss: 0.7720
```

```
Epoch 3/15
```

```
50000/50000 [=====] - 1s 16us/step - loss: 0.6263 - acc: 0.8544 - val_loss: 0.6263
```

```
Epoch 4/15
```

```
50000/50000 [=====] - 1s 16us/step - loss: 0.5564 - acc: 0.8647 - val_loss: 0.5564
```

```
Epoch 5/15
```

```
50000/50000 [=====] - 1s 19us/step - loss: 0.5141 - acc: 0.8711 - val_loss: 0.5141
```

```
Epoch 6/15
```

```
50000/50000 [=====] - 1s 17us/step - loss: 0.4850 - acc: 0.8762 - val_loss: 0.4850
```

```
Epoch 7/15
```

```

50000/50000 [=====] - 1s 16us/step - loss: 0.4636 - acc: 0.8798 - val_l
Epoch 8/15
50000/50000 [=====] - 1s 17us/step - loss: 0.4470 - acc: 0.8829 - val_l
Epoch 9/15
50000/50000 [=====] - 1s 17us/step - loss: 0.4336 - acc: 0.8855 - val_l
Epoch 10/15
50000/50000 [=====] - 1s 16us/step - loss: 0.4225 - acc: 0.8875 - val_l
Epoch 11/15
50000/50000 [=====] - 1s 16us/step - loss: 0.4131 - acc: 0.8893 - val_l
Epoch 12/15
50000/50000 [=====] - 1s 16us/step - loss: 0.4051 - acc: 0.8917 - val_l
Epoch 13/15
50000/50000 [=====] - 1s 16us/step - loss: 0.3980 - acc: 0.8928 - val_l
Epoch 14/15
50000/50000 [=====] - 1s 17us/step - loss: 0.3918 - acc: 0.8941 - val_l
Epoch 15/15
50000/50000 [=====] - 1s 16us/step - loss: 0.3863 - acc: 0.8954 - val_l

```

```

In [11]: print('Test score:', score[0])
         print('Test accuracy:', score[1])

```

```

('Test score:', 0.3645026022195816)

```

```

('Test accuracy:', 0.9024)

```

2 comparing the validation and test accuracies for different set of Learning rate and Batch Size

Learning Rate	Val Acc	Test Acc
0.001	0.9144	0.9197
0.01	0.9141	0.9210
0.05	1.9153	0.9216
0.1	0.9170	0.9214

Batch Size	Val Acc	Test Acc
1	0.9139	0.9207
32	0.9206	0.9261
128	0.9208	0.926
1024	0.9208	0.926